

Title:Uber & Lyft Cab prices Analysis

```
In [ ]: Name: Allu Pranaya
RollNo: 2211CS010016
Section: S1-04
```

The cab_rides dataset contains 2,093 records and 10 columns representing cab ride information. It includes ride distance, cab type, source and destination locations, price, surge multiplier, ride ID, product ID, and cab name. The numeric columns (distance, price, surge_multiplier, time_stamp) allow analysis of ride trends and relationships, such as how distance affects price or how surge pricing varies. Categorical columns (cab_type, source, destination, product_id, name) help in grouping rides to find the most popular routes, cab types, or destinations. Some observations: Price has missing values (1920 non-null), which may need handling before analysis. Distance ranges across different rides, useful for fare vs. distance analysis. Surge multiplier shows how peak demand affects pricing. The dataset can be used to explore ride patterns, fare distributions, peak demand effects, and the most popular cab types and locations. This dataset is suitable for data visualization, correlation analysis, and pattern discovery in cab rides.

```
In [1]: sc
```

Out[1]: **SparkContext**

[Spark UI](#)

Version	v3.5.6
Master	local[*]
AppName	PySparkShell

```
In [5]: from pyspark.sql import SparkSession
from pyspark.sql.functions import col, corr
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Initialize Spark
spark = SparkSession.builder.appName("CabRidesAnalysis").getOrCreate()
```

```
In [6]: df = spark.read.csv("cab_rides.csv", header=True, inferSchema=True)
df.show(5) # Preview first 5 rows
df.printSchema() # Check schema
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|distance|cab_type|time_stamp| destination| source|price|surge_multiplier| id| product_id
| name|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
| 0.44| Lyft|1.54495E12|North Station|Haymarket Square| 5.0| 1.0|424553bb-7174-41e...| lyft_line
| Shared|
| 0.44| Lyft|1.54328E12|North Station|Haymarket Square| 11.0| 1.0|4bd23055-6827-41c...| lyft_premier
| Lux|
| 0.44| Lyft|1.54337E12|North Station|Haymarket Square| 7.0| 1.0|981a3613-77af-462...| lyft
| Lyft|
| 0.44| Lyft|1.54355E12|North Station|Haymarket Square| 26.0| 1.0|c2d88af2-d278-4bf...| lyft_luxsuv
|Lux Black XL|
| 0.44| Lyft|1.54346E12|North Station|Haymarket Square| 9.0| 1.0|e0126e1f-8ca9-4f2...| lyft_plus
| Lyft XL|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 5 rows

```

```

root
|-- distance: double (nullable = true)
|-- cab_type: string (nullable = true)
|-- time_stamp: double (nullable = true)
|-- destination: string (nullable = true)
|-- source: string (nullable = true)
|-- price: double (nullable = true)
|-- surge_multiplier: double (nullable = true)
|-- id: string (nullable = true)
|-- product_id: string (nullable = true)
|-- name: string (nullable = true)

```

```

In [7]: # Summary statistics
df.describe().show()

```

summary	distance	cab_type	time_stamp	destination	source	price	surge_multiplier
	id	product_id	name				
count	2093	2093	2093	2093	2093	1920	2093
	2093	2093	2093	2093			
mean	2.1312326803631096	NULL	1.544018408982322E12	NULL	NULL	16.398177083333334	1.010630673674152
	NULL	NULL	NULL	NULL	NULL	9.083548960021902	0.08311274424308637
stddev	1.0509020571359042	NULL	6.86041100059527E8	NULL	NULL	9.083548960021902	0.08311274424308637
	NULL	NULL	NULL	NULL	NULL	9.083548960021902	0.08311274424308637
min	0.3	Lyft	1.54321E12	Back Bay	Back Bay	3.0	1.0
	001e4903-c2ec-475...	55c66225-fbe7-4fd...	Black				
max	7.46	Uber	1.54516E12	West End	West End	67.5	2.5
	ffe072dd-eb41-496...	lyft_premier	WAV				

```
In [8]: numeric_cols = ['distance', 'time_stamp', 'price', 'surge_multiplier']

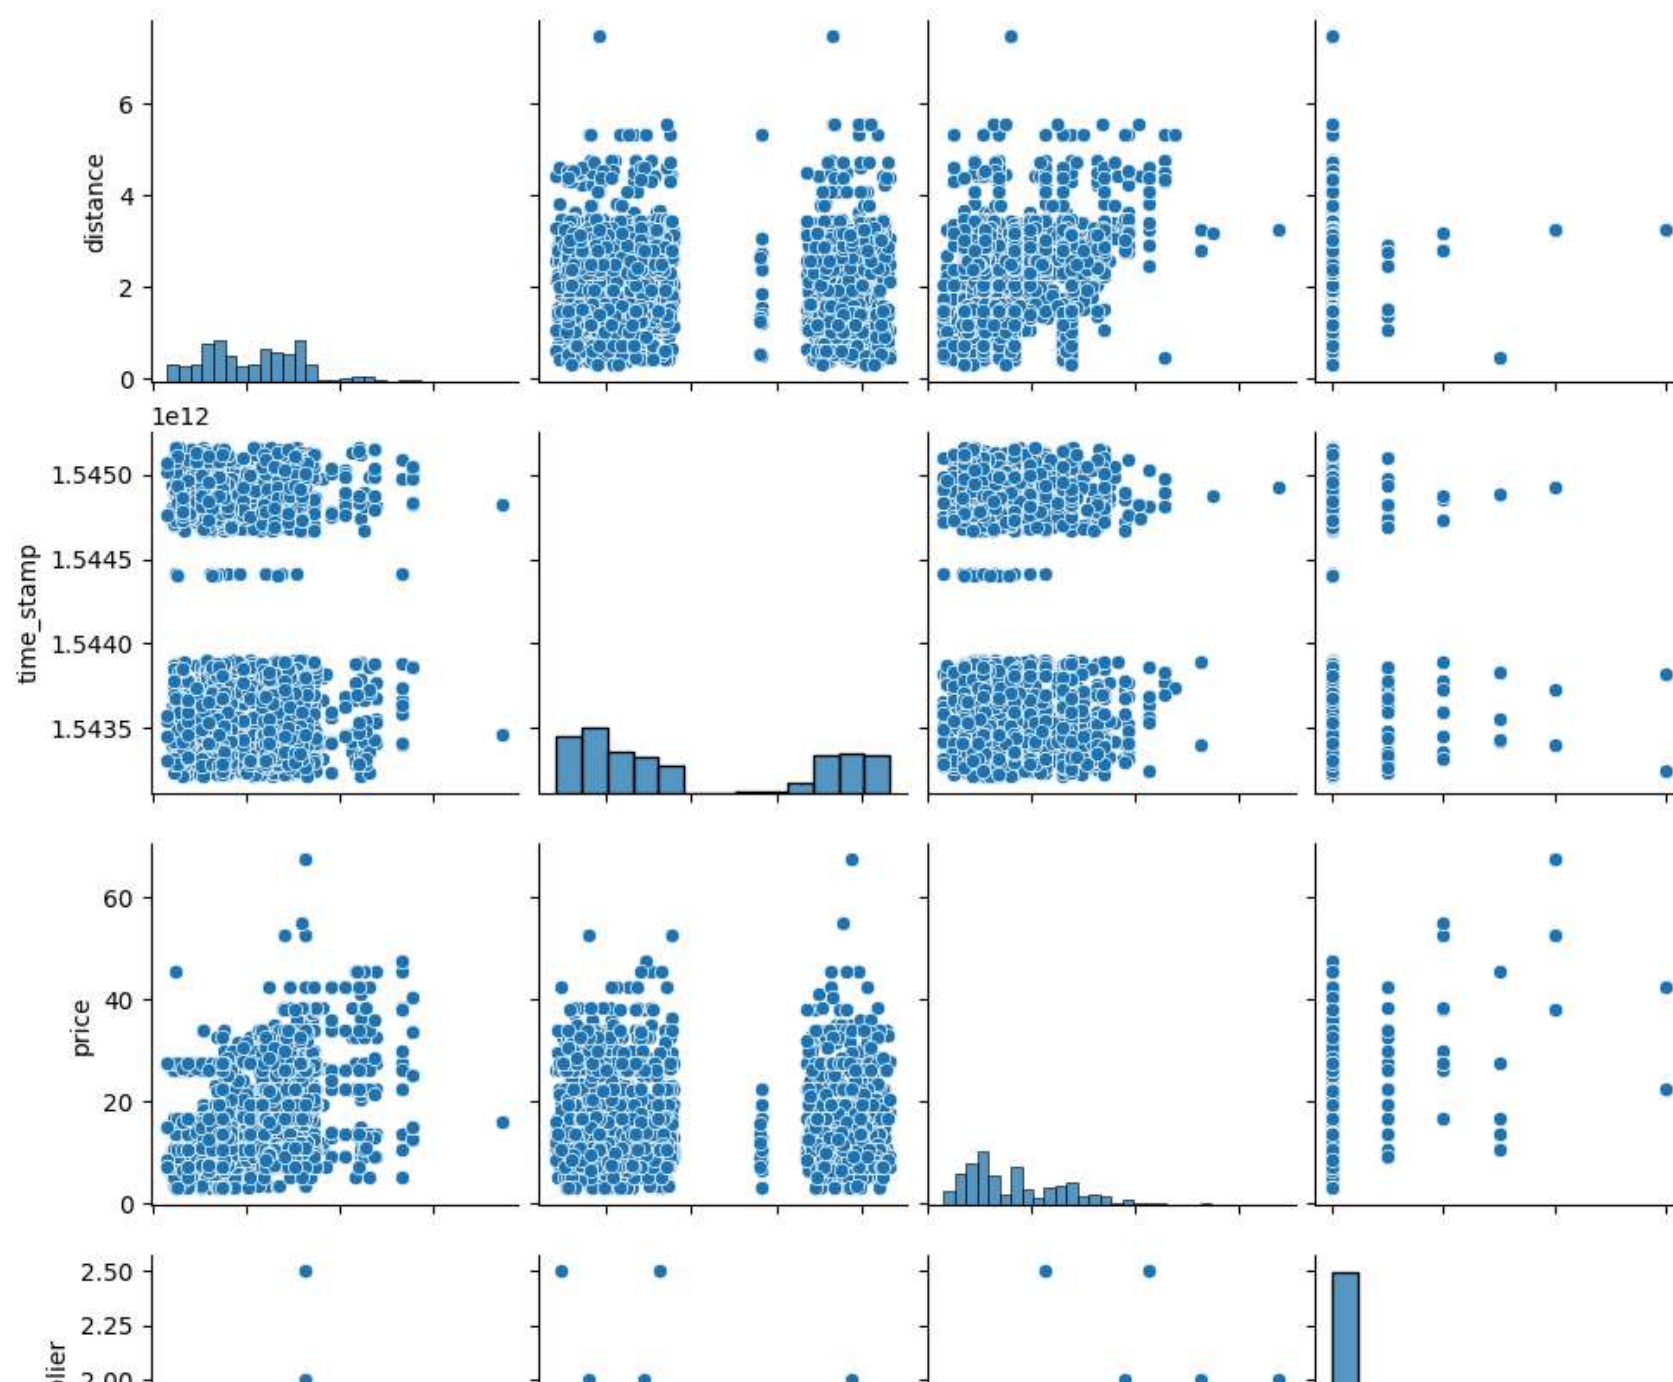
# Print correlations
for col1 in numeric_cols:
    for col2 in numeric_cols:
        if col1 != col2:
            corr_value = df.stat.corr(col1, col2)
            print(f"Correlation between {col1} and {col2}: {corr_value:.2f}")
```


```
Correlation between distance and time_stamp: 0.02
Correlation between distance and price: 0.28
Correlation between distance and surge_multiplier: 0.02
Correlation between time_stamp and distance: 0.02
Correlation between time_stamp and price: -0.00
Correlation between time_stamp and surge_multiplier: -0.02
Correlation between price and distance: 0.28
Correlation between price and time_stamp: -0.00
Correlation between price and surge_multiplier: 0.20
Correlation between surge_multiplier and distance: 0.02
Correlation between surge_multiplier and time_stamp: -0.02
Correlation between surge_multiplier and price: 0.20
```

```
In [9]: # Convert PySpark DataFrame to Pandas  
pdf = df.toPandas()
```

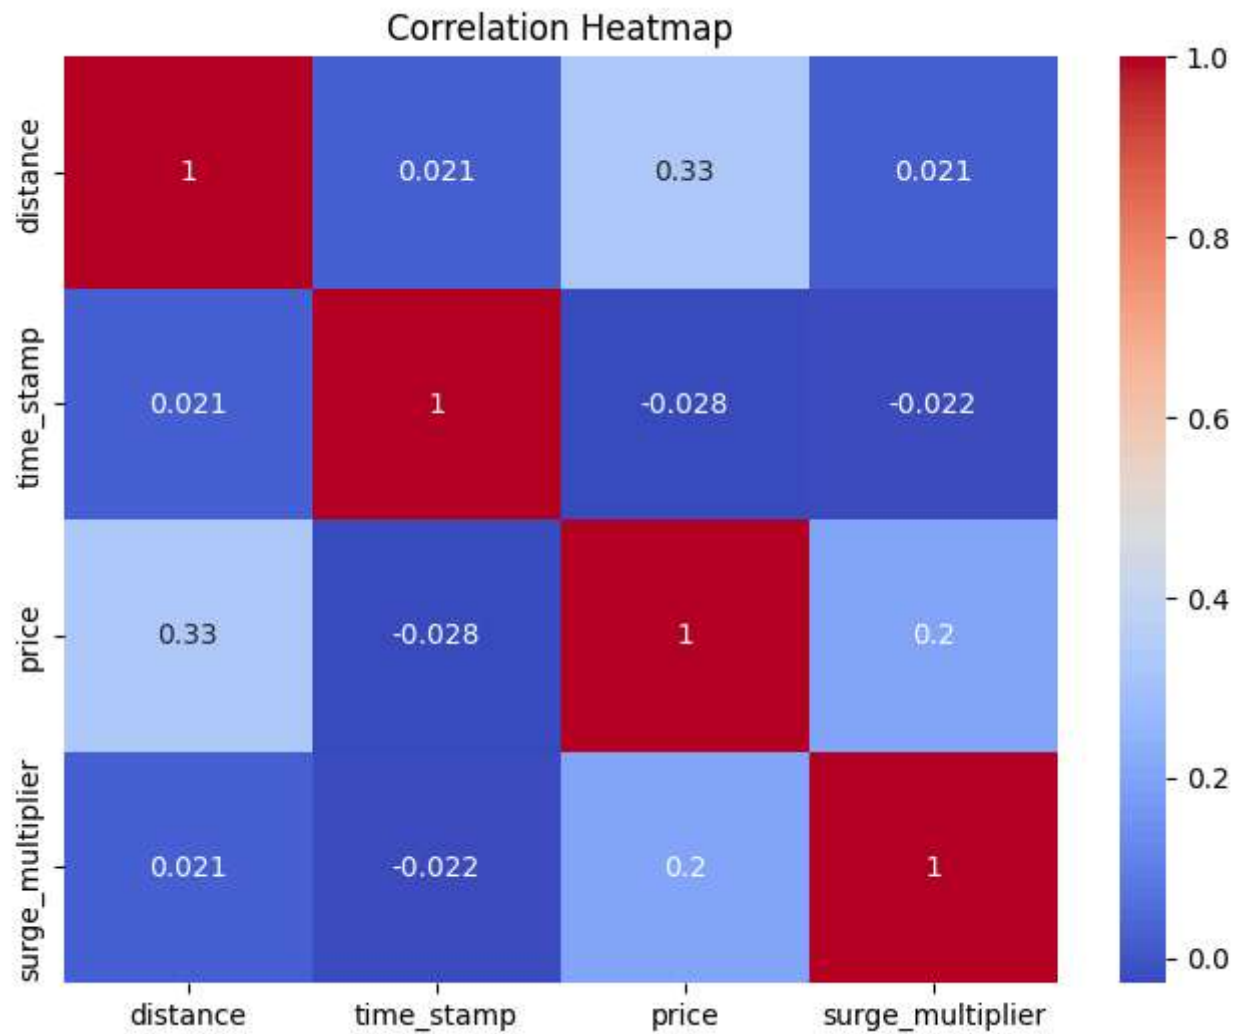
```
In [10]: sns.pairplot(pdf[numeric_cols])  
plt.suptitle("Pairplot of Numeric Columns", y=1.02)  
plt.show()
```

Pairplot of Numeric Columns

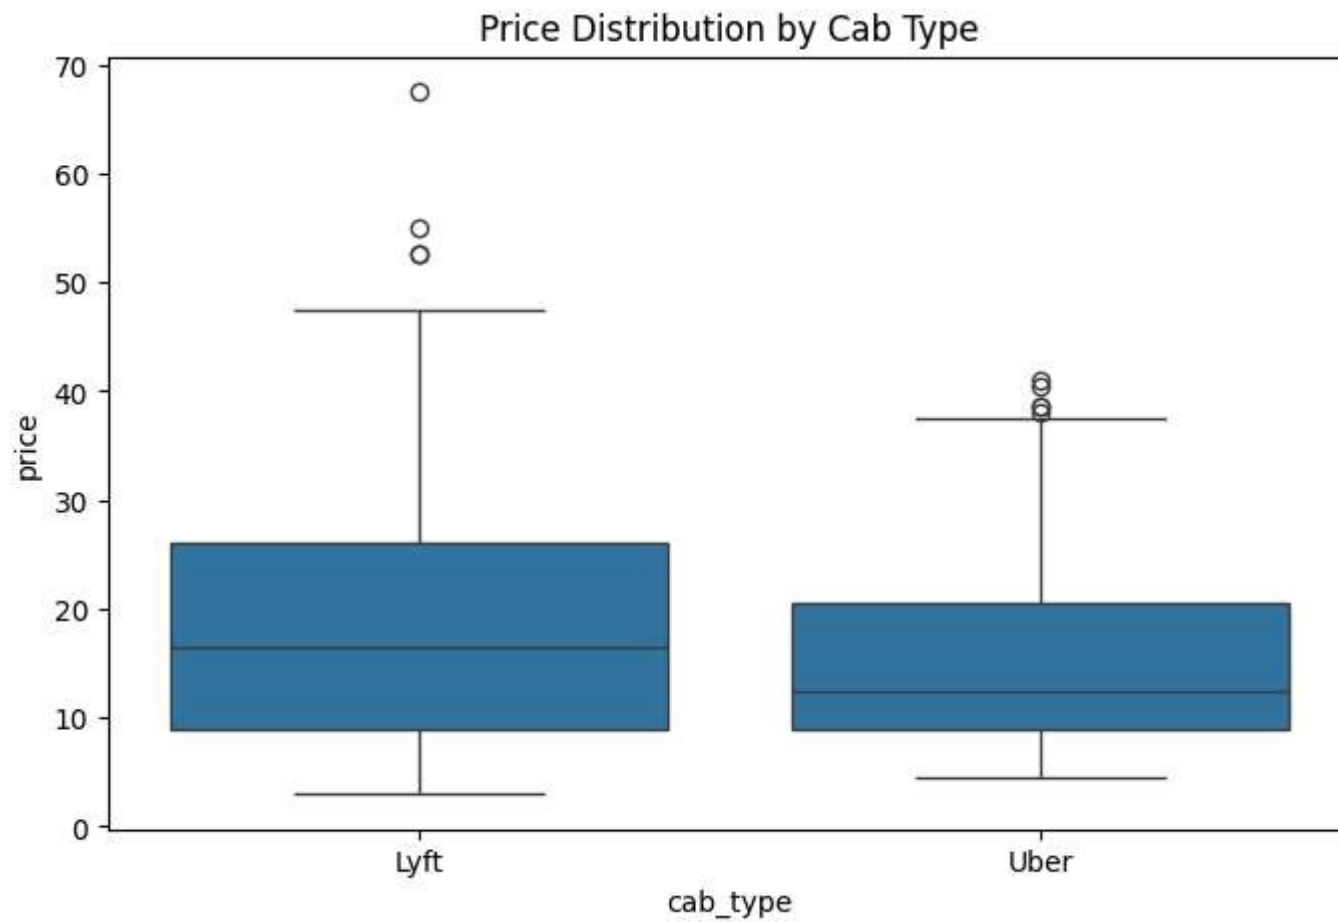




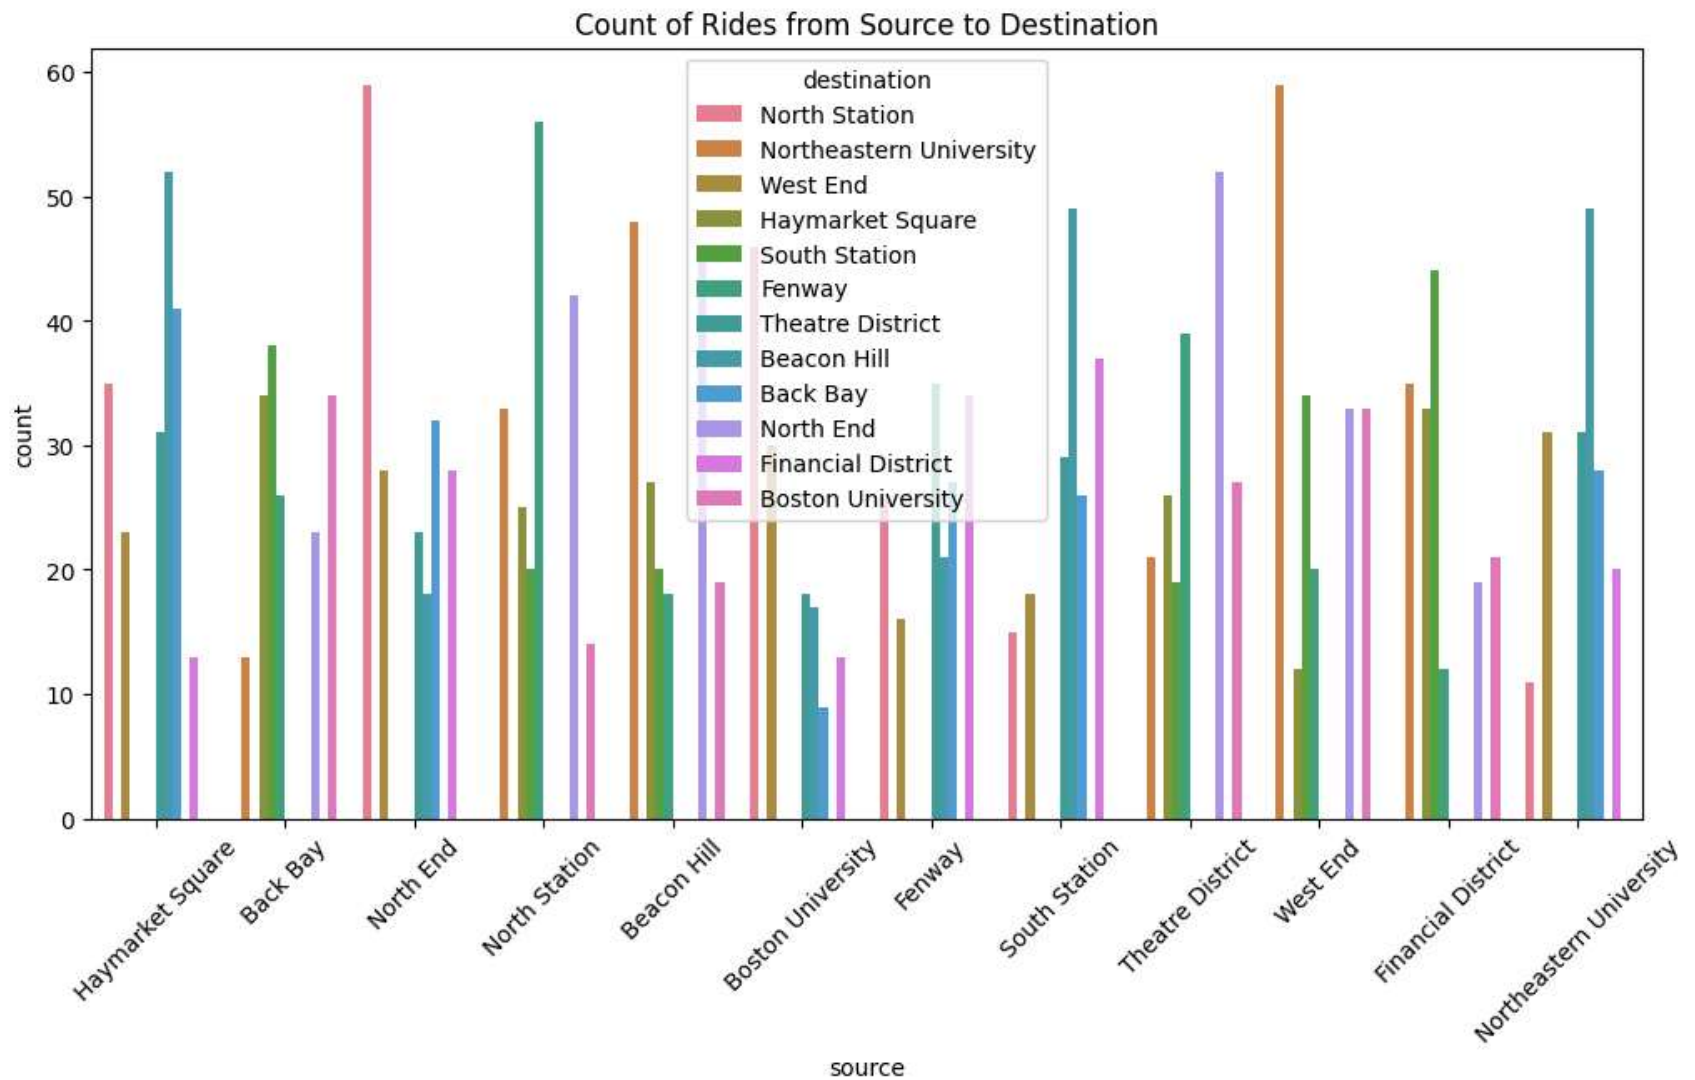
```
In [11]: corr_matrix = pdf[numeric_cols].corr()
plt.figure(figsize=(8,6))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```



```
In [12]: plt.figure(figsize=(8,5))
sns.boxplot(x='cab_type', y='price', data=pdf)
plt.title("Price Distribution by Cab Type")
plt.show()
```



```
In [13]: plt.figure(figsize=(12,6))
sns.countplot(x='source', hue='destination', data=pdf)
plt.xticks(rotation=45)
plt.title("Count of Rides from Source to Destination")
plt.show()
```

```
In [14]: # Average price by cab type
avg_price = pdf.groupby('cab_type')['price'].mean()
print("Average Price by Cab Type:")
print(avg_price)

# Most popular source
popular_source = pdf['source'].value_counts().idxmax()
print(f"Most popular source location: {popular_source}")
```

Average Price by Cab Type:

cab_type

Lyft 17.322598

Uber 15.554781

Name: price, dtype: float64

Most popular source location: Haymarket Square

```
In [16]: from pyspark.sql import SparkSession
        from pyspark.sql.functions import col, avg, max, min, count
        import matplotlib.pyplot as plt
        import seaborn as sns
        import pandas as pd

        spark = SparkSession.builder.appName("CabRidesAnalysis2").getOrCreate()
```

```
In [17]: df = spark.read.csv("cab_rides.csv", header=True, inferSchema=True)
        df.show(5)
        df.printSchema()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|distance|cab_type|time_stamp| destination| source|price|surge_multiplier| id| product_id
| name|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
| 0.44| Lyft|1.54495E12|North Station|Haymarket Square| 5.0| 1.0|424553bb-7174-41e...| lyft_line
| Shared|
| 0.44| Lyft|1.54328E12|North Station|Haymarket Square| 11.0| 1.0|4bd23055-6827-41c...| lyft_premier
| Lux|
| 0.44| Lyft|1.54337E12|North Station|Haymarket Square| 7.0| 1.0|981a3613-77af-462...| lyft
| Lyft|
| 0.44| Lyft|1.54355E12|North Station|Haymarket Square| 26.0| 1.0|c2d88af2-d278-4bf...| lyft_luxsuv
|Lux Black XL|
| 0.44| Lyft|1.54346E12|North Station|Haymarket Square| 9.0| 1.0|e0126e1f-8ca9-4f2...| lyft_plus
| Lyft XL|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 5 rows

```

```

root
|-- distance: double (nullable = true)
|-- cab_type: string (nullable = true)
|-- time_stamp: double (nullable = true)
|-- destination: string (nullable = true)
|-- source: string (nullable = true)
|-- price: double (nullable = true)
|-- surge_multiplier: double (nullable = true)
|-- id: string (nullable = true)
|-- product_id: string (nullable = true)
|-- name: string (nullable = true)

```

```

In [18]: avg_price_cab = df.groupBy("cab_type").agg(avg("price").alias("avg_price"))
avg_price_cab.show()

```

```
+-----+-----+
| cab_type | avg_price |
+-----+-----+
| Lyft | 17.32259825327511 |
| Uber | 15.554780876494023 |
+-----+-----+
```

```
In [19]: max_distance_cab = df.groupBy("cab_type").agg(max("distance").alias("max_distance"))
max_distance_cab.show()
```

```
+-----+-----+
| cab_type | max_distance |
+-----+-----+
| Lyft | 5.33 |
| Uber | 7.46 |
+-----+-----+
```

```
In [20]: rides_per_source = df.groupBy("source").agg(count("*").alias("num_rides"))
rides_per_source.show()
```


```
+-----+-----+
| source | num_rides |
+-----+-----+
| Financial District | 164 |
| Northeastern Univ... | 170 |
| North End | 188 |
| Boston University | 133 |
| North Station | 190 |
| Back Bay | 168 |
| Theatre District | 184 |
| South Station | 174 |
| Fenway | 159 |
| Haymarket Square | 195 |
| West End | 191 |
| Beacon Hill | 177 |
+-----+-----+
```

```
In [21]: avg_surge_destination = df.groupBy("destination").agg(avg("surge_multiplier").alias("avg_surge"))
avg_surge_destination.show()
```

destination	avg_surge
Financial District	1.0
Northeastern Univ...	1.0
North End	1.0
Boston University	1.008445945945946
North Station	1.0
Back Bay	1.0368098159509203
Theatre District	1.0224550898203593
South Station	1.0214285714285714
Fenway	1.0146198830409356
Haymarket Square	1.0079617834394905
West End	1.0256849315068493
Beacon Hill	1.0

```
In [22]: pdf_avg_price = avg_price_cab.toPandas()
pdf_max_distance = max_distance_cab.toPandas()
pdf_rides_source = rides_per_source.toPandas()
pdf_avg_surge = avg_surge_destination.toPandas()
```

```
In [23]: plt.figure(figsize=(6,4))
sns.barplot(x='cab_type', y='avg_price', data=pdf_avg_price)
plt.title("Average Price per Cab Type")
plt.show()
```

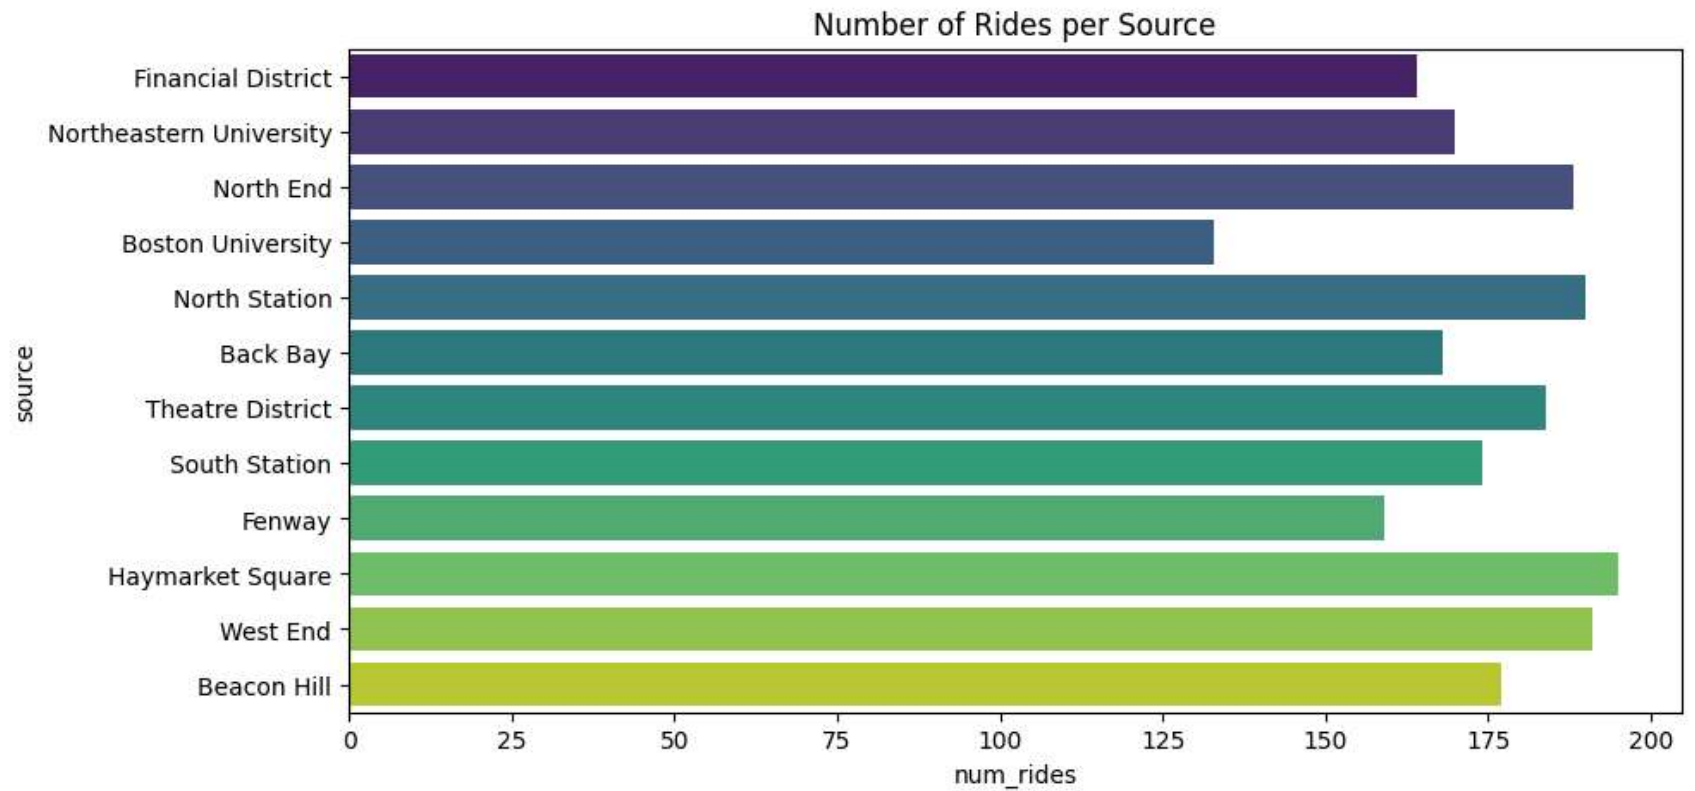


```
In [24]: plt.figure(figsize=(10,5))
sns.barplot(x='num_rides', y='source', data=pdf_rides_source, palette="viridis")
plt.title("Number of Rides per Source")
plt.show()
```


C:\Users\Harini\AppData\Local\Temp\ipykernel_9356\2514970801.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='num_rides', y='source', data=pdf_rides_source, palette="viridis")
```



```
In [25]: plt.figure(figsize=(6,5))
sns.scatterplot(x='distance', y='price', hue='cab_type', data=df.toPandas())
plt.title("Price vs Distance by Cab Type")
plt.show()
```



```
In [27]: # Most expensive cab type on average
expensive_cab = pdf_avg_price.loc[pdf_avg_price['avg_price'].idxmax(), 'cab_type']
print(f"The most expensive cab type on average is: {expensive_cab}")

# Source with highest number of rides
top_source = pdf_rides_source.loc[pdf_rides_source['num_rides'].idxmax(), 'source']
print(f"The source with highest number of rides is: {top_source}")

# Destination with highest average surge
top_surge_destination = pdf_avg_surge.loc[pdf_avg_surge['avg_surge'].idxmax(), 'destination']
print(f"The destination with highest average surge multiplier is: {top_surge_destination}")
```


The most expensive cab type on average is: Lyft

The source with highest number of rides is: Haymarket Square

The destination with highest average surge multiplier is: Back Bay

In []: